



# **BIOS226 - Topic 5 - Supervised Learning (Part 2)**

## **The Tumor Subtype Classification Pipeline**

Dr. Robert Treharne

# From Data to Decision: The Pipeline

We will use our synthetic breast cancer gene expression dataset generated in R.

- 120 patients
- 10,000 gene features
- Subtypes: Luminal\_A vs Basal\_like
- Logistic regression classifier

Goal: Predict tumor subtype from gene expression profiles.

# 1. Define the Prediction Problem

Clinical question: Can we classify a tumor as *Basal\_like* or *Luminal\_A* based on gene expression?

Model task: Binary classification.

- Input: 10,000 gene expression values per patient
- Output: Probability tumor is *Basal\_like*

Positive class: **Basal\_like**

## 2. Validate the Data Structure

Before modeling, verify:

- Column structure ( Patient\_ID , Subtype , Gene\_1 ... Gene\_n )
- Exactly two classes present
- All gene columns are numeric
- No schema inconsistencies

Why? Garbage in leads to garbage out.

### 3. Stratified Train/Test Split

Split the data:

- 80% training
- 20% test
- Stratified by subtype

Why stratified? It maintains class balance in both sets.

The test set is held out and **never** used during training.

## 4. Feature Selection (Train Only)

Rank genes by absolute difference in mean expression between classes.

Select top K genes (e.g., 25).

Important: Feature selection is performed on training data only.

Why? To avoid data leakage.

## 5. Scaling & Normalisation (Train Only)

Gene expression values vary in magnitude.

Standardise each selected gene:

- Subtract mean
- Divide by standard deviation

Scaling parameters are learned from training data only and then applied to test data.

## 6. Cross-Validation on Training Data

Perform k-fold cross-validation (e.g., 5-fold):

1. Split training set into 5 parts
2. Train on 4 folds
3. Validate on 1 fold
4. Repeat 5 times

Outputs:

- AUC per fold
- Precision per fold

Purpose: Estimate **model stability** before touching test set.

## 7. Fit the Logistic Regression Model

Logistic regression estimates:

$P(\text{Basal\_like} \mid \text{gene expression})$

Model form:

$$\log(p / (1 - p)) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots$$

**Output:** Probability between 0 and 1 for each patient.

## 8. Confusion Matrix: Understanding Errors

Using chosen threshold (e.g., 0.85 - high specificity! Why?), compute:

- True Positives (TP)
- False Positives (FP)
- True Negatives (TN)
- False Negatives (FN)


From this we calculate:

- Precision
- Sensitivity
- Specificity

This describes performance at one threshold.

**Truth Table (Confusion Matrix)**

Precision=0.750 Sensitivity=0.900 Specificity=0.800



## 9. ROC Curve & AUC

The model outputs probabilities.

If threshold varies, sensitivity and false positive rate change.


ROC curve shows performance across all thresholds:

- X-axis: False Positive Rate
- Y-axis: True Positive Rate

AUC measures discrimination ability:

- 0.5 = random
- 1.0 = perfect

ROC Curves: CV Mean +/- SD Envelope + Test (Test AUC = 0.863)



# 10. Final Test Evaluation and Discipline

After all training steps, evaluate once on the held-out test set.

Report:

- Confusion matrix
- Precision
- ROC AUC
- MSE (probability-based)
- R squared (probability-based)

This gives an unbiased estimate of real-world performance.

# Pipeline Summary

1. Define question
2. Validate data
3. Split train/test
4. Select features (train only)
5. Scale data (train only)
6. Cross-validate
7. Fit model
8. Evaluate with confusion matrix
9. Evaluate with ROC/AUC
10. Report final test performance